

# Introduction to Java Database Connectivity API

CSCE 156 - Introduction to Computer Science II

Christopher M. Bourke  
cbourke@cse.unl.edu

# Java Database Connectivity API

## Java Database Connectivity (JDBC)

- ▶ General API (using interfaces) for Java client code to connect/interact with a database
- ▶ Database providers (IBM, Oracle, etc.) provide *drivers*
- ▶ Driver: specific implementation of the API for interacting with a particular database
- ▶ Support for
  - ▶ `Connection`
  - ▶ `PreparedStatement`
  - ▶ `ResultSet`
  - ▶ Common Java data types ( `Integer`, `Double`, `String` )

## JDBC: basic step-by-step

1. Load the database JDBC driver  
Note: your particular driver (.jar file) must be in the class or build path of your project
2. Make a connection to the database
3. Formulate your query(ies) & prepare your statement (set parameters)
4. Execute your query
5. If its a `SELECT` query:
  - 5.1 Get your *result set*
  - 5.2 Process your results
6. Clean up your resources (close resources, close connection)

# JDBC

## Reflectively loading a driver

- ▶ For portability, applications written toward JDBC API, not a particular driver
- ▶ Driver is loaded at run time through reflection
- ▶ Could be made configurable or delegated by some controller

```
1  try {  
2      Class.forName("com.mysql.jdbc.Driver").newInstance();  
3  } catch (InstantiationException e) {  
4      ...  
5  } catch (IllegalAccessException e) {  
6      ...  
7  } catch (ClassNotFoundException e) {  
8      ...  
9  }
```

# JDBC

## Connection

Java provides connectivity through `java.sql.Connection`:

```
1 String url = "jdbc:mysql://cse.unl.edu/cselogin";
2 String u = "cselogin";
3 String p = "mysqlpasswd";
4 Connection conn = null;
5 try {
6     conn = DriverManager.getConnection(url, u, p);
7 } catch (SQLException sqle) {
8     ...
9 }
```

# JDBC

## Transactions

- ▶ By default, all queries are auto-commit
- ▶ To change this, use `conn.setAutoCommit(false)`
- ▶ No changes committed until `conn.commit()` is called
- ▶ Implicitly: new transaction after each commit
- ▶ Able to explicitly rollback using `conn.rollback()`
- ▶ Some drivers may also support `conn.setReadOnly(true)`

# JDBC I

## Querying – Prepared Statement

- ▶ Always good to use `PreparedStatement`
- ▶ Can define *parameters* using `?`
- ▶ Parameters indexed by `1..n`
- ▶ Can be reused (parameters reset and required)
- ▶ Parameters are *safe*!
- ▶ Special characters are *escaped*
- ▶ Potentially unsafe SQL code is *sanitized*

# JDBC II

## Querying – Prepared Statement

```
1 String query = "SELECT last_name AS lastName " +
2               "FROM user WHERE nuid = ?";
3 PreparedStatement ps = null;
4 try {
5     ps = conn.prepareStatement(query);
6     ps.setString(1, "35140602");
7 } catch (SQLException sqle) {
8     ...
9 }
```



# JDBC I

## Querying – Result Sets

- ▶ `executeQuery()` is for read-only (select statements)
- ▶ Select statements return *results*: columns and rows
- ▶ Results are encapsulated in a Java `ResultSet` object
- ▶ Initially a result set “points” just before the first *row*
- ▶ Iterating through a `ResultSet` : `rs.next()`
- ▶ Returns a boolean: true if the iteration was successful, false otherwise
- ▶ If successful, the “current” result row is now pointed to
- ▶ Columns can be referenced by name (or alias) using a `String` or
- ▶ Columns can be accessed via index (1,...,)
- ▶ Standard getters provide functionality to get-and-cast columns

# JDBC II

## Querying – Result Sets

```
1  ResultSet rs = null;
2  try {
3      rs = ps.executeQuery();
4      while(rs.next()) {
5          Integer nuid = rs.getInt("nuid")
6          String firstName = rs.getString("firstName");
7      }
8  } catch (SQLException sqle) {
9      ...
10 }
```

# JDBC

## Querying – Updates

- ▶ Always use a prepared statement!
- ▶ Same syntax holds for INSERT statements

```
1 String query = "UPDATE user SET email = ?, " +
2               "last_updated = ? WHERE nuid = ?";
3 PreparedStatement ps = null;
4 try {
5     ps = conn.prepareStatement(query);
6     ps.setString(1, "cmbourke@gmail.com");
7     ps.setString(2, "2011-01-01 00:00:01");
8     ps.setString(3, "35140602");
9     ps.executeUpdate();
10 } catch (SQLException sqle) {
11     ...
12 }
```

# JDBC

## Good Practices – Rethrow Exceptions

- ▶ Most methods explicitly throw `SQLException`
- ▶ This is a *checked exception* that *must* be caught and handled
- ▶ Occurs with DB errors or program bugs
- ▶ Little can be done either way
- ▶ Good to catch, log and rethrow
- ▶ Even better: use a logging utility like log4j

```
1  ...
2  } catch (SQLException sqle) {
3      System.out.println("SQLException: ");
4      sqle.printStackTrace();
5      throw new RuntimeException(sqle);
6  }
```

# JDBC

## Cleaning Up

- ▶ Objects hold onto valuable external resources
- ▶ Network traffic (keep alive), limited connection pool, etc.
- ▶ Best practice to release resources as soon as they are no longer needed: `close()` method

```
1  try {
2      if(rs != null && !rs.isClosed())
3          rs.close();
4      if(ps != null && !ps.isClosed())
5          ps.close();
6      if(conn != null && !conn.isClosed())
7          conn.close();
8  } catch (SQLException e) {
9      ...
10 }
```

# JDBC

## Full Example Demonstration

A full demonstration is available in the `unl.cse.jdbc` package in the SVN.

Demonstration based on the student/course enrollment database.

# Good Practice Tip 1

## ALWAYS use Prepared Statements

When available, in any framework or language, always use prepared statements

- ▶ Safer
- ▶ Better for batch queries
- ▶ Myth: no performance hit
- ▶ Protects against *injection attacks*
- ▶ Using just one method: more uniform, less of a chance of a mistake
- ▶ Unfortunately: some frameworks support *named parameters*, not JDBC

# Injection Attack

## Example

- ▶ Say we pull a string value from a web form ( `lastName` )
- ▶ Not using a prepared statement:

```
1 String query = "SELECT primary_email FROM user " +  
2               "WHERE last_name = '"+lastName+"'";
```

- ▶ Without scrubbing the input, say a user enters:

```
a';DROP TABLE user;
```

- ▶ Actual query run:

```
SELECT primary_email FROM user WHERE last = 'a';DROP TABLE users;
```

- ▶ Another example: input: `"' OR '1'='1'"`

- ▶ Actual query:


```
SELECT primary_email FROM user WHERE last_name = '' OR '1'='1'
```

- ▶ In detail: <https://www.netsparker.com/blog/web-security/sql-injection-cheat-sheet/>



# Injection Attack

## Example



# Good Practice Tip 2

## Enumerate fields in SELECT statements

- ▶ Using `SELECT * ...` grabs all fields even if you don't use them
- ▶ Be intentional about what you want/need, only the minimal set
- ▶ Allows the database to optimize, reduces network traffic
- ▶ Protects against table changes
- ▶ Use aliasing ( `first_name AS firstName` ) on all fields to reduce affects of changes to field names

# Additional Issues

## Additional Issues

- ▶ Security Issues
  - ▶ Where/how should database passwords be stored?
  - ▶ Good security policy: *assume* an attacker has your password & take the necessary precautions (secure the server and network)
  - ▶ Do not store sensitive data unencrypted
- ▶ Efficiency Issues
  - ▶ **Repeat:** **close your resources**
  - ▶ Connection Pools
  - ▶ Good normalization, design, & practice

# Resources

- ▶ MySQL 5.1 Reference Manual  
(<http://dev.mysql.com/doc/refman/5.1/en/index.html>)
- ▶ MySQL Community Server (<http://www.mysql.com/downloads/>)
- ▶ MySQL Workbench – a MySQL GUI (<http://wb.mysql.com/>)
- ▶ Connector/J – MySQL JDBC connector  
(<http://www.mysql.com/downloads/connector/j/>)
- ▶ Stanford's *Introduction to Databases* free online course:  
<http://db-class.com/>

# Log4j I

- ▶ Home: <http://logging.apache.org/>
- ▶ Standard output is not appropriate for deployed applications (may not exist, certainly no one is “listening” to it)
- ▶ Logging provides valuable debugging, metrics, and auditing information
- ▶ Provides runtime configuration, formatting, rolling files, etc.

# Log4j II

- ▶ Supports granularity in different levels of logging ( `ERROR` to `DEBUG` )
- ▶ Usage: give each loggable class a static logger:

```
1  private static org.apache.log4j.Logger log =  
2      Logger.getLogger(MyClass.class);
```

- ▶ Then use it: `log.error("SQLException: ", e);`
- ▶ Configure using a `log4j.properties` file (must be in the class path)
- ▶ Or: call `BasicConfigurator.configure();` in your `main` to have a default setup

# Exercise

*Write basic CRUD methods for the `Employee` / `Person` tables by writing static methods to insert, delete, retrieve and update records in both tables.*